

eda

October 28, 2023

0.1 Data Exploration and Preprocessing

Importing Necessary Libraries

```
[122]: import pandas as pd
import numpy as np
import missingno
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
import warnings
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler
warnings.filterwarnings('ignore')
```

Reading the Data

```
[123]: data = pd.read_csv("bank.csv")
data.head()
```

```
[123]:
```

	age	job	marital	education	default	balance	housing	loan	contact	\
0	59	admin.	married	secondary	no	2343	yes	no	unknown	
1	56	admin.	married	secondary	no	45	no	no	unknown	
2	41	technician	married	secondary	no	1270	yes	no	unknown	
3	55	services	married	secondary	no	2476	yes	no	unknown	
4	54	admin.	married	tertiary	no	184	no	no	unknown	

	day	month	duration	campaign	pdays	previous	poutcome	deposit
0	5	may	1042	1	-1	0	unknown	yes
1	5	may	1467	1	-1	0	unknown	yes
2	5	may	1389	1	-1	0	unknown	yes
3	5	may	579	1	-1	0	unknown	yes
4	5	may	673	2	-1	0	unknown	yes


EXPLORATORY DATA ANALYSIS:

```
[124]: print("Shape of the dataset: ", data.shape)
```

Shape of the dataset: (11162, 17)

```
[125]: missingno.matrix(data)
```

```
[125]: <Axes: >
```



```
[126]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 11162 entries, 0 to 11161
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         11162 non-null  int64
1   job         11162 non-null  object
2   marital     11162 non-null  object
3   education   11162 non-null  object
4   default     11162 non-null  object
5   balance     11162 non-null  int64
6   housing     11162 non-null  object
7   loan        11162 non-null  object
8   contact     11162 non-null  object
9   day         11162 non-null  int64
10  month       11162 non-null  object
11  duration    11162 non-null  int64
12  campaign    11162 non-null  int64
13  pdays       11162 non-null  int64
14  previous    11162 non-null  int64
15  poutcome    11162 non-null  object
16  deposit     11162 non-null  object
dtypes: int64(7), object(10)
memory usage: 1.4+ MB
```

```
[127]: data.describe()
```

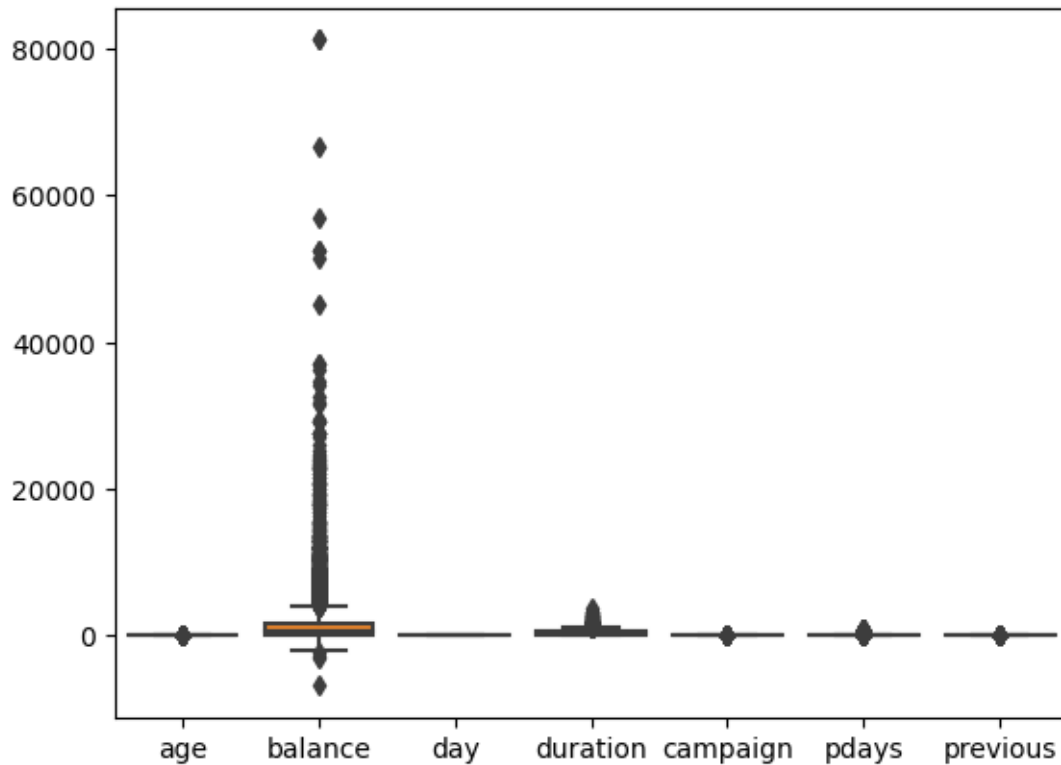
```
[127]:
```

	age	balance	day	duration	campaign \
count	11162.000000	11162.000000	11162.000000	11162.000000	11162.000000
mean	41.231948	1528.538524	15.658036	371.993818	2.508421
std	11.913369	3225.413326	8.420740	347.128386	2.722077
min	18.000000	-6847.000000	1.000000	2.000000	1.000000
25%	32.000000	122.000000	8.000000	138.000000	1.000000
50%	39.000000	550.000000	15.000000	255.000000	2.000000
75%	49.000000	1708.000000	22.000000	496.000000	3.000000
max	95.000000	81204.000000	31.000000	3881.000000	63.000000

	pdays	previous
count	11162.000000	11162.000000
mean	51.330407	0.832557
std	108.758282	2.292007
min	-1.000000	0.000000
25%	-1.000000	0.000000
50%	-1.000000	0.000000
75%	20.750000	1.000000
max	854.000000	58.000000

```
[128]: sns.boxplot(data=data)
```

```
[128]: <Axes: >
```



AGE:

```
[129]: data['age'].info()
```

```
<class 'pandas.core.series.Series'>
RangeIndex: 11162 entries, 0 to 11161
Series name: age
Non-Null Count  Dtype
-----
11162 non-null  int64
dtypes: int64(1)
memory usage: 87.3 KB
```


```
[130]: data['age'].value_counts().head(10)
```

```
[130]: age
31    496
32    477
34    466
33    464
35    461
30    456
36    432
```

```
37    370
38    353
39    343
Name: count, dtype: int64
```

```
[131]: sns.histplot(x='age', data=data, color='green')
```

```
[131]: <Axes: xlabel='age', ylabel='Count'>
```



JOB:

```
[132]: data['job'].describe()
```


```
[132]: count          11162
unique           12
top      management
freq           2566
Name: job, dtype: object
```

```
[133]: data['job'].value_counts()
```

```
[133]: job
      management      2566
      blue-collar    1944
      technician     1823
      admin.         1334
      services       923
      retired        778
      self-employed  405
      student        360
      unemployed     357
      entrepreneur   328
      housemaid      274
      unknown        70
      Name: count, dtype: int64
```

```
[134]: px.pie(data, values=np.ones(11162), names='job',
      ↪title='job',color_discrete_sequence=px.colors.sequential.Brwnyl)
```


```
[135]: sns.countplot(x='job', data=data, hue='deposit')
      plt.xticks(rotation=-65)
      plt.show()
```



MARITAL:

```
[136]: sns.countplot(x="marital", data=data)
```


```
[136]: <Axes: xlabel='marital', ylabel='count'>
```



```
[137]: fig = px.pie(data, values=np.ones(11162), names='marital',  
    ↪title='marital',color_discrete_sequence=px.colors.sequential.Brwnyl)  
fig.show()
```


```
[138]: sns.countplot(x="marital", data=data, hue="deposit")
```

```
[138]: <Axes: xlabel='marital', ylabel='count'>
```





EDUCATION:

```
[139]: sns.countplot(x="education", data=data, color='rosybrown')  
plt.xticks(rotation=-65)  
plt.show()
```



```
[140]: sns.countplot(x="education", data=data, hue="deposit")  
plt.xticks(rotation=-65)  
plt.show()
```




```
[141]: fig = px.pie(data, values=np.ones(11162), names='education',
    ↪title='education',color_discrete_sequence=px.colors.sequential.Brwnyl)
fig.show()
```

DEFAULT:


```
[142]: sns.countplot(x="default", data=data)
```

```
[142]: <Axes: xlabel='default', ylabel='count'>
```



```
[143]: sns.countplot(x="default", data=data, hue="deposit")
```

```
[143]: <Axes: xlabel='default', ylabel='count'>
```




```
[144]: fig = px.pie(data, values=np.ones(11162), names='default',  
               title='default',color_discrete_sequence=px.colors.sequential.Brwnyl)  
fig.show()
```

BALANCE:


```
[145]: sns.boxplot(y="balance",data=data)
```

```
[145]: <Axes: ylabel='balance'>
```



```
[146]: sns.kdeplot(x="balance",data =data,hue="deposit")
```

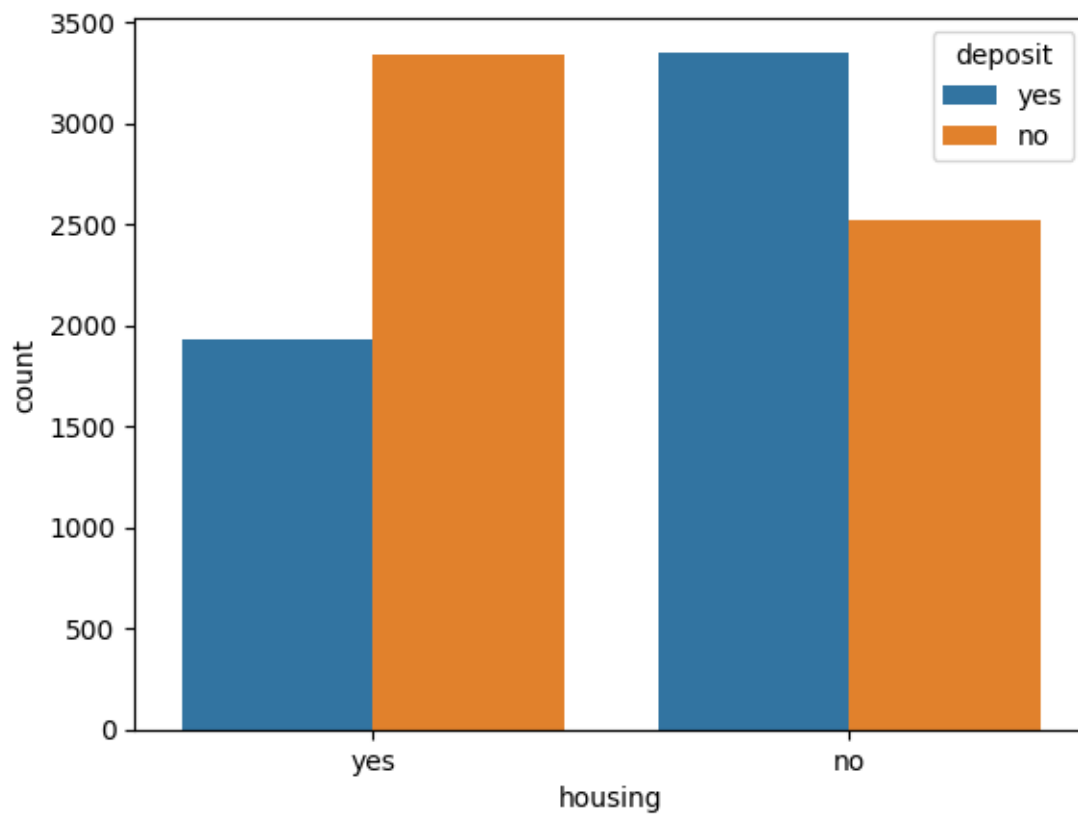
```
[146]: <Axes: xlabel='balance', ylabel='Density'>
```



HOUSING:

```
[147]: sns.countplot(x="housing", data=data, hue="deposit")
```


```
[147]: <Axes: xlabel='housing', ylabel='count'>
```



LOAN:

```
[148]: sns.countplot(x="loan", data=data, hue="deposit")
```


```
[148]: <Axes: xlabel='loan', ylabel='count'>
```

DURATION:

```
[149]: sns.kdeplot(x="duration",data =data,hue ="deposit")
```


```
[149]: <Axes: xlabel='duration', ylabel='Density'>
```



CAMPAIGN:

```
[150]: sns.kdeplot(data=data,x="campaign",hue="deposit")
```

```
[150]: <Axes: xlabel='campaign', ylabel='Density'>
```




POUTCOME:

```
[151]: fig = px.pie(data, values=np.ones(11162), names='poutcome',
    ↪title='poutcome',color_discrete_sequence=px.colors.sequential.Brwnyl)
fig.show()
```

```
[152]: sns.countplot(x="poutcome", data=data, hue="deposit")
```

```
[152]: <Axes: xlabel='poutcome', ylabel='count'>
```



DEPOSIT:

```
[153]: fig = px.pie(data, values=np.ones(11162), names='deposit',
    ↳ title='deposit', color_discrete_sequence=px.colors.sequential.Brwnyl)
    fig.show()
```

CORRELATION:

```
[154]: num = data.select_dtypes('int64')
    num
```

```
[154]:
```


	age	balance	day	duration	campaign	pdays	previous
0	59	2343	5	1042	1	-1	0
1	56	45	5	1467	1	-1	0
2	41	1270	5	1389	1	-1	0
3	55	2476	5	579	1	-1	0
4	54	184	5	673	2	-1	0
...
11157	33	1	20	257	1	-1	0
11158	39	733	16	83	4	-1	0
11159	32	29	19	156	2	-1	0
11160	43	0	8	9	2	172	5

```
11161    34         0     9      628         1     -1         0
```

```
[11162 rows x 7 columns]
```

```
[155]: sns.heatmap(num.corr(),annot =True)
```

```
[155]: <Axes: >
```



DATA PREPROCESSING:

LABEL ENCODING:

```
[156]: data.head()
```

```
[156]:   age      job  marital  education  default  balance  housing  loan  contact  \
0   59  admin.  married  secondary     no      2343     yes   no  unknown
1   56  admin.  married  secondary     no         45     no   no  unknown
2   41 technician  married  secondary     no      1270     yes   no  unknown
3   55  services  married  secondary     no      2476     yes   no  unknown
4   54  admin.  married  tertiary     no        184     no   no  unknown
```

	day	month	duration	campaign	pdays	previous	poutcome	deposit
0	5	may	1042	1	-1	0	unknown	yes
1	5	may	1467	1	-1	0	unknown	yes
2	5	may	1389	1	-1	0	unknown	yes
3	5	may	579	1	-1	0	unknown	yes
4	5	may	673	2	-1	0	unknown	yes

```
[157]: from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler

le = LabelEncoder()

for i in data.select_dtypes('object').columns:
    data[i] = le.fit_transform(data[i])
```

```
[158]: data
```

```
[158]:
```

	age	job	marital	education	default	balance	housing	loan	contact	\
0	59	0	1	1	0	2343	1	0	2	
1	56	0	1	1	0	45	0	0	2	
2	41	9	1	1	0	1270	1	0	2	
3	55	7	1	1	0	2476	1	0	2	
4	54	0	1	2	0	184	0	0	2	
...	
11157	33	1	2	0	0	1	1	0	0	
11158	39	7	1	1	0	733	0	0	2	
11159	32	9	2	1	0	29	0	0	0	
11160	43	9	1	1	0	0	0	1	0	
11161	34	9	1	1	0	0	0	0	0	

	day	month	duration	campaign	pdays	previous	poutcome	deposit
0	5	8	1042	1	-1	0	3	1
1	5	8	1467	1	-1	0	3	1
2	5	8	1389	1	-1	0	3	1
3	5	8	579	1	-1	0	3	1
4	5	8	673	2	-1	0	3	1
...
11157	20	0	257	1	-1	0	3	0
11158	16	6	83	4	-1	0	3	0
11159	19	1	156	2	-1	0	3	0
11160	8	8	9	2	172	5	0	0
11161	9	5	628	1	-1	0	3	0

[11162 rows x 17 columns]

```
[159]: data.columns
```

```
[159]: Index(['age', 'job', 'marital', 'education', 'default', 'balance', 'housing',
            'loan', 'contact', 'day', 'month', 'duration', 'campaign', 'pdays',
            'previous', 'poutcome', 'deposit'],
            dtype='object')
```

```
[167]: x = data.drop(['contact', 'day', 'month', 'pdays', 'previous', 'deposit'], axis = 1)
x
```

```
[167]:
```

	age	job	marital	education	default	balance	housing	loan	\
0	59	0	1	1	0	2343	1	0	
1	56	0	1	1	0	45	0	0	
2	41	9	1	1	0	1270	1	0	
3	55	7	1	1	0	2476	1	0	
4	54	0	1	2	0	184	0	0	
...	
11157	33	1	2	0	0	1	1	0	
11158	39	7	1	1	0	733	0	0	
11159	32	9	2	1	0	29	0	0	
11160	43	9	1	1	0	0	0	1	
11161	34	9	1	1	0	0	0	0	

	duration	campaign	poutcome
0	1042	1	3
1	1467	1	3
2	1389	1	3
3	579	1	3
4	673	2	3
...
11157	257	1	3
11158	83	4	3
11159	156	2	3
11160	9	2	0
11161	628	1	3

[11162 rows x 11 columns]

```
[168]: y=data["deposit"]
```

CUSTOMER SEGMENTATION AND MODELLING:


```
[172]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(x, y, test_size = 0.2,
↳ random_state = 32)
```

```
[169]: from sklearn.cluster import KMeans
from sklearn.preprocessing import MinMaxScaler
```

```
ms = MinMaxScaler()

X_scaled = ms.fit_transform(x)
```

```
[173]: cs = []
for i in range(1, 11):
    kmeans = KMeans(n_clusters = i, init = 'k-means++', max_iter = 300, n_init=
    ↪ 10, random_state = 0)
    kmeans.fit(X_train)
    cs.append(kmeans.inertia_)
plt.plot(range(1, 11), cs,marker='o')
plt.title('The Elbow Method')
plt.xlabel('Number of clusters')
plt.ylabel('CS')
plt.show()
```



```
[174]: kmeans = KMeans(n_clusters = 2, init = 'k-means++', max_iter = 300, n_init =
    ↪ 10, random_state = 0)
kmeans.fit(x)
```



```
[174]: KMeans(n_clusters=2, n_init=10, random_state=0)
```

```
[177]: correct_labels = sum(y== kmeans.labels_)
print("Result: %d out of %d samples were correctly labeled." % (correct_labels,
    ↪y.size))
print('Accuracy score: {0:0.2f}'. format(correct_labels/float(y.size)))
```

Result: 5221 out of 11162 samples were correctly labeled.
Accuracy score: 0.47

```
[178]: # Using PCA:

from sklearn.preprocessing import StandardScaler, normalize
from sklearn.decomposition import PCA
```

```
[179]: # Standardize data
scaler = StandardScaler()
scaled_df = scaler.fit_transform(x)

# Normalizing the Data
normalized_df = normalize(scaled_df)


# Converting the numpy array into a pandas DataFrame
normalized_df = pd.DataFrame(normalized_df)

# Reducing the dimensions of the data
pca = PCA(n_components = 2)
X_principal = pca.fit_transform(normalized_df)
X_principal = pd.DataFrame(X_principal)
X_principal.columns = ['P1', 'P2']

X_principal.head(2)
```

```
[179]:          P1          P2
0  0.608737  0.173193
1  0.343215 -0.189492
```

```
[180]: sse = {} #sum of sqaured errors
for k in range(1, 10):
    kmeans = KMeans(n_clusters=k, max_iter=1000).fit(X_principal)
    sse[k] = kmeans.inertia_ # Inertia: Sum of distances of samples to their
    ↪closest cluster center
plt.figure()
plt.plot(list(sse.keys()), list(sse.values()),marker='o')
plt.xlabel("Number of cluster")
plt.ylabel("SSE")
plt.show()
```




```
[181]: # Silhoutte coefficient method:
```

```
from sklearn.metrics import silhouette_score
```

```
[182]: silhouette_scores = []
```

```
for n_cluster in range(2, 10):  
    silhouette_scores.append(  
        silhouette_score(X_principal, KMeans(n_clusters = n_cluster).  
            ↪fit_predict(X_principal)))
```

```
# Plotting a bar graph to compare the results  
k = [2, 3, 4, 5, 6, 7, 8, 9]  
plt.bar(k, silhouette_scores)  
plt.xlabel('Number of clusters', fontsize = 10)  
plt.ylabel('Silhouette Score', fontsize = 10)  
plt.show()
```




```
[183]: # k=3:

kmeans = KMeans(n_clusters=3)
kmeans.fit(X_principal)
```

```
[183]: KMeans(n_clusters=3)
```

```
[184]: # Visualizing the clustering
plt.scatter(X_principal['P1'], X_principal['P2'],
            c = KMeans(n_clusters = 3).fit_predict(X_principal), cmap =plt.cm.
            ↪winter)
plt.show()
```



```
[185]: # Hierarchical Agglomerative Clustering:

from sklearn.cluster import AgglomerativeClustering

ac = AgglomerativeClustering(n_clusters=2, linkage='ward',
    ↪compute_full_tree=False)
model_ac = ac.fit(X_train)
train_predicted = model_ac.fit_predict(X_train)
test_predicted = model_ac.fit_predict(X_test)
print(train_predicted.shape, y_train.shape)
print(test_predicted.shape, y_test.shape)

(8929,) (8929,)
(2233,) (2233,)
```

```
[187]: # Accuracy:

from sklearn.metrics import accuracy_score
train_accuracy = accuracy_score(y_train, train_predicted)
test_accuracy = accuracy_score(y_test, test_predicted)

from sklearn.metrics import classification_report
```

```
print("Train: ",classification_report(y_train,train_predicted))
print("Test: ", classification_report(y_test,test_predicted))
```

Train:		precision	recall	f1-score	support
	0	0.53	0.99	0.69	4687
	1	0.56	0.01	0.02	4242
	accuracy			0.53	8929
	macro avg	0.54	0.50	0.35	8929
	weighted avg	0.54	0.53	0.37	8929

Test:		precision	recall	f1-score	support
	0	0.41	0.03	0.06	1186
	1	0.46	0.95	0.62	1047
	accuracy			0.46	2233
	macro avg	0.44	0.49	0.34	2233
	weighted avg	0.44	0.46	0.32	2233

```
[189]: # DBSCAN:

from sklearn.cluster import DBSCAN

dbs = DBSCAN(eps=3, min_samples=500)

model_dbs = dbs.fit(X_train)
predicted_train = list(map(lambda x: x+1, model_dbs.fit_predict(X_train)))
predicted_test = list(map(lambda x: x+1, model_dbs.fit_predict(X_test)))
```

```
[190]: from sklearn.metrics import accuracy_score
train_accuracy = accuracy_score(y_train, predicted_train)
test_accuracy = accuracy_score(y_test, predicted_test)

print("Train: \n",classification_report(y_train,predicted_train))
print("Test: \n",classification_report(y_test,predicted_test))
```

Train:		precision	recall	f1-score	support
	0	0.52	1.00	0.69	4687
	1	0.00	0.00	0.00	4242
	accuracy			0.52	8929

macro avg	0.26	0.50	0.34	8929
weighted avg	0.28	0.52	0.36	8929

Test:

	precision	recall	f1-score	support
0	0.53	1.00	0.69	1186
1	0.00	0.00	0.00	1047
accuracy			0.53	2233
macro avg	0.27	0.50	0.35	2233
weighted avg	0.28	0.53	0.37	2233

[]: